E.L.L.A.

Embedded Local Logic Agent

User Manual

v3.9.0

© 2026 Andre Zabel, Berlin · ella-agent.de

E.L.L.A. — User Manual

Version 1.2.0

Hello — I'm E.L.L.A.

E.L.L.A. stands for *Embedded Local Logic Agent* — your personal AI assistant for Windows. Always ready, always local, always yours.

What makes me different: I run entirely on your own machine. Your data stays with you — no cloud server, no subscription, no privacy fine print. Just install and go.

Getting Started

Installation

- Run `E.L.L.A Setup 1.2.0.exe`
- Choose an installation path (default: `C:\Program Files\E.L.L.A`)
- The installer sets everything up automatically — MariaDB, Ollama, Node.js. Nothing else to install.

First Launch

On first launch, ELLA guides you through a short onboarding flow:

- **Language selection** — German or English
- **EULA** — accept the terms of use and the ELLA Directive
- **Hardware scan** — ELLA detects your system and recommends a suitable language model
- **Model selection** — you choose the model that fits your hardware (no download without your consent)
- **Voice selection** — 4 female voices in DE/EN/ES/FR with preview
- **Location** — for weather and local information (optional)
- **License activation** — enter your license key `ELLA-XXXX-XXXX-XXXX-XXXX`
- **Chat** — you're ready

The Chat

How to Talk to Me

Just write what you need — in plain language, no commands, no special syntax. I understand natural language.

- "How much disk space is left?"
- "Write an email to Klaus saying I can't make it tomorrow"
- "What's on the calendar this week?"
- "Open Spotify"

Voice Input

Click the **microphone button** in the input bar — or simply say **"Ella"** if the wakeword is enabled. ELLA listens, converts your speech to text, and responds.

Wakeword modes:

- **Off** — no automatic listening
- **Manual** — microphone button only
- **Wakeword** — listen for "Ella", then record
- **Always** — continuous listening

Voice Output (TTS)

ELLA responds with voice too. Two engines are available:

- **Kokoro TTS** (default) — local neural voice, fully offline, natural sound
- **Microsoft Edge TTS** (fallback) — requires internet, very reliable

The voice language can be set independently of the app language: German, English, Spanish, or French.

Feedback

Below each response you'll find **thumbs up / thumbs down**. When you give feedback, ELLA learns from it and adjusts her responses to your preferences over time.

Tool Display

When ELLA does something — searches a file, opens an app, reads your calendar — you briefly see which tool she used. This gives you insight into what's happening.

What I Can Do For You

PC & System

I know Windows inside out. Ask me anything about your machine.

Just say:

- "Open Chrome"
- "How much RAM am I using right now?"
- "Show me all running processes"
- "Kill the process XY"
- "Set the volume to 40 percent"
- "What's in my clipboard?"
- "Shut down / restart the computer"
- "Which programs start automatically with Windows?"

Drives & Storage

Just say:

- "How much space is left on drive C?"
- "Find duplicate files on my drives"
- "Format my USB stick" *(confirmation required)*
- "Rename drive E to Data"
- "Empty the recycle bin"

Files & Documents

I can find, open, and create files across all your drives.

Just say:

- "Find the file called Proposal"
- "Open the Projects folder"
- "Create a new Word document"
- "What does this PDF say?" *(attach file or name the path)*
- "Read me the Excel spreadsheet"
- "Summarize this document"

Dictation & Writing

Just say:

- "Dictate a letter to Ms. Miller"
- "Write this in the style of my emails"

- "Correct this text"
- "Translate this to German"

Mail & Calendar

Outlook:

- "Show my unread emails"
- "Mark the email as read / spam / archived"
- "Draft an email to Klaus: meeting tomorrow at 2pm?"
- "What's on my calendar this week?"
- "Create an appointment: dentist, Friday at 10am"

Gmail & Google Calendar (OAuth required, set up once in Settings):

- "Show me my Gmail emails"
- "What's on my Google Calendar today?"

Tasks

ELLA maintains a persistent task list. Tasks survive restarts.

Just say:

- "Add the task 'File taxes'"
- "Show my open tasks"
- "Mark 'File taxes' as done"
- "Delete task XY"

Alarms

Just say:

- "Set an alarm for 7:30"
- "Set an alarm for tomorrow morning at 6"
- "Which alarms are active?"
- "Delete the 7:30 alarm"

Alarms are checked daily and trigger a notification plus a TTS announcement.

Reminders

Reminders are set via the Windows Task Scheduler and appear as popups at the specified time.

Just say:

- "Remind me tomorrow at 3pm: call the doctor's office"
- "What reminders do I have?"
- "Delete reminder XY"

Music & Entertainment

Just say:

- "Play something by Coldplay"
- "Open my media player with the album 'Brothers in Arms'"
- "Find Metallica tracks in my library"

Network & Monitoring

Just say:

- "Who's on my network?"
- "Show active TCP connections"
- "What's my Wi-Fi status?"
- "Show me the firewall rules"

Screenshot & Screen Analysis

(Requires OpenAI key)

Just say:

- "Take a screenshot and analyze what you see"
- "Describe what's happening on screen"
- "Analyze this image" *(paste an image)*

Multi-Agents

The Multi-Agent feature is ELLA's most powerful mode. Multiple specialized AI agents work **in parallel**, delivering more in less time than a single agent ever could.

> ***Note:*** *Multi-Agents requires a configured OpenAI key (Settings → API Key). Ollama is not used for this feature.*

How to Access It

Navigate to **Multi-Agents** in the sidebar.

Predefined Pipelines

■ Morning Briefing

Three agents work simultaneously:

- Agent 1 reads today's calendar appointments
- Agent 2 retrieves your open tasks
- Agent 3 fetches the current weather for your location

The result: a compact daily briefing combining all three.

■■ System Report

Four agents analyze simultaneously:

- Agent 1: Free space on all drives
- Agent 2: RAM and CPU usage
- Agent 3: Running processes (top 10 by resource consumption)
- Agent 4: Network status and open connections

The result: a complete system health report.

Live View

While the pipeline runs, you see a live card for each agent showing:

- Real-time status (running / done / error)
- Live-streamed agent text
- Which tools the agent is currently using

Run a Setup After Completion *(optional)*

You can link a pipeline with a **Setup**. After all agents finish, ELLA automatically executes the setup — for example, opening or closing specific apps.

Select the desired setup from the dropdown before starting the pipeline.

The Data Is Real

What you see in Multi-Agent view is not sample data — it's real information from your system, your calendar, and your drives. Just like in the regular chat, only faster and parallelized.

Setups

Setups are macros — define once which apps to open and close, then trigger them with a single phrase.

Just say:

- "Load my work setup"
- "Start the gaming setup"

Create setups: **Settings** → **Setups** → **New**.

Memory

ELLA remembers information about you — preferences, habits, notes. Memory is separated by profile and persists permanently.

Just say:

- "Remember: I don't like meetings before 9am"
- "What do you know about me?"
- "Forget the note about XY"

Proactive Messages

ELLA reaches out proactively when something important happens:

- **Drive nearly full** — when a drive drops below 10% free space
- **High RAM usage** — under sustained load
- **USB device connected** — detection and announcement of new drives
- **Alarm triggered** — voice alert at the set time
- **Clipboard pattern** — IBAN, tracking number, URL, phone number detected → ELLA asks if she can help
- **Meeting preparation** — 15 minutes before a calendar appointment
- **End of day** — evening summary of open tasks
- **Break reminder** — after a long work session

User Profiles

ELLA supports multiple user profiles on the same machine. Each profile has:

- Its own memory

- Its own conversation history
- Its own optional PIN protection

Manage profiles: **Settings** → **Profiles**.

Homework Mode

Homework Mode locks the computer to a defined app whitelist. Only approved programs can be opened — ideal for children while studying.

Important: Only the correct PIN ends the mode. There is no automatic expiry time — this is intentional, so children cannot simply wait it out.

Activate: Say "Activate homework mode" or go to **Settings** → **Homework Mode**.

Accessibility

Blind Mode

Enable under **Settings** → **Accessibility** → **Blind Mode**:

- TTS always active — every response is read aloud
- Responses without Markdown formatting (better for screen readers)
- Audio tones as status feedback (loading, error, success)
- Welcome announcement on startup
- All buttons and inputs labeled with `aria-label`

Deaf Mode

Enable under **Settings** → **Accessibility** → **Deaf Mode**:

- TTS and microphone button completely hidden
- Visual banner in chat
- No mention of voice features in system prompt

Settings

Open Settings via the sidebar or say "Open settings".

Key Areas

Area	What you'll find
General	Language, name, location
LLM Provider	Ollama (local) or OpenAI — or Auto
API Key	OpenAI key for GPT features and Multi-Agents
Voice	TTS engine, language, stress analysis
Tools	Which tools ELLA is allowed to use
Apps	List of your installed programs
Setups	Macros for program combinations
Rules	Allow / deny individual actions
Profiles	Manage user profiles
Homework	Configure children's mode
Accessibility	Blind and deaf mode

Rules

You can fine-tune ELLA's behavior with rules — no coding required:

- "Create a rule: ELLA must not send emails"
- "Create a rule: ELLA must not set alarms between 10pm and 7am"

Rules can be time-based and are enforced by the rule engine before any tool runs.

The ELLA Directive

ELLA has four absolute prohibitions — hardcoded into the architecture, not overridable by any setting or rule:

- **harm** — ELLA never causes harm: no physical, financial, psychological, or data-related damage
- **conceal** — ELLA never conceals: no hidden actions, no false statements about capabilities
- **surveil** — ELLA never monitors without awareness and explicit consent
- **exfiltrate** — ELLA never shares data without explicit consent

When ELLA declines a request, you'll see a red **DIRECTIVE** banner in the chat with the reason.

Privacy

- All data stays on your machine
- No telemetry, no tracking, no analytics
- The AI model (Ollama) runs fully offline
- OpenAI features (screenshot analysis, Multi-Agents) send data to OpenAI — only when you've configured a key and actively use the feature
- OpenAI does not train models on API data

Frequently Asked Questions

ELLA isn't responding — what do I do?

Check if Ollama is running. The orange banner at the top of the chat appears when Ollama is unreachable. Restart Ollama or switch to OpenAI mode.

Can I use ELLA offline?

Yes — completely. All core features run without internet. Only weather (Open-Meteo), Gmail/Google Calendar (OAuth), and OpenAI features need a connection.

How do I change the voice?

Settings → Voice → Language. Choose between German, English, Spanish, and French. The app language is not affected.

Can ELLA support multiple users on one machine?

Yes — through user profiles. Each profile has its own memory and history. Optionally protect with a PIN.

What's the difference between Chat and Multi-Agents?

In Chat, a single agent (Ollama or OpenAI) works sequentially. In Multi-Agents, multiple specialized OpenAI agents work simultaneously — faster and with more depth for complex queries.

ELLA declined something — why?

Either one of your own rules is active, or the ELLA Directive applies. The red banner explains the reason. Rules can be adjusted in Settings.

How do I clear the memory?

Say "Forget everything about me" — or delete individual entries with "Show me your memory" followed by "Forget entry number X".

E.L.L.A. v1.2.0 — Embedded Local Logic Agent

Built for Windows 10/11 (x64)